

Redge regression

```
In [1]: import numpy as np  
import pandas as pd
```

reading csv file

```
In [2]: kcd=pd.read_csv("/home/placement/Downloads/fiat500.csv")
kcd.info
```

```
Out[2]: <bound method DataFrame.info of
wners \
0      1 lounge      51      882  25000      1
1      2 pop        51     1186  32500      1
2      3 sport      74     4658 142228      1
3      4 lounge     51     2739 160000      1
4      5 pop        73     3074 106880      1
...    ...
1533 1534 sport      51     3712 115280      1
1534 1535 lounge     74     3835 112000      1
1535 1536 pop        51     2223  60457      1
1536 1537 lounge     51     2557  80750      1
1537 1538 pop        51     1766  54276      1
```

```
lat lon price
0 44.907242 8.611560 8900
1 45.666359 12.241890 8800
2 45.503300 11.417840 4200
3 40.633171 17.634609 6000
4 41.903221 12.495650 5700
...
1533 45.069679 7.704920 5200
1534 45.845692 8.666870 4600
1535 45.481541 9.413480 7500
1536 45.000702 7.682270 5990
1537 40.323410 17.568270 7900
```

```
[1538 rows x 9 columns]>
```

In [3]: kcd

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [4]: a=kcd.groupby(['model']).count()
a

Out[4]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
model								
lounge	1094	1094	1094	1094	1094	1094	1094	1094
pop	358	358	358	358	358	358	358	358
sport	86	86	86	86	86	86	86	86

removing id,lat,lon columns

```
In [5]: drop=kcd.drop(['ID', 'lat', 'lon'],axis=1)  
drop
```

Out[5]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

converting strings into integers

```
In [6]: drop['model']=drop['model'].map({'lounge':1,'pop':2,'sport':3})
drop
```

Out[6]:

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [7]: cor1=drop.corr()
```

importing seaborn and matplotlib

```
In [8]: import seaborn as sb
import matplotlib.pyplot as mp
```

extracting price column

```
In [9]: y=drop['price']  
x=drop.drop('price',axis=1)  
x
```

Out[9]:

	model	engine_power	age_in_days	km	previous_owners
0	1	51	882	25000	1
1	2	51	1186	32500	1
2	3	74	4658	142228	1
3	1	51	2739	160000	1
4	2	73	3074	106880	1
...
1533	3	51	3712	115280	1
1534	1	74	3835	112000	1
1535	2	51	2223	60457	1
1536	1	51	2557	80750	1
1537	2	51	1766	54276	1

1538 rows × 5 columns

In [10]: y

Out[10]:

0	8900
1	8800
2	4200
3	6000
4	5700

...

1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

test and train

```
In [11]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [12]: x_train.head(10)
```

```
Out[12]:
```

	model	engine_power	age_in_days	km	previous_owners
527	1	51	425	13111	1
129	1	51	1127	21400	1
602	2	51	2039	57039	1
331	1	51	1155	40700	1
323	1	51	425	16783	1
1358	1	51	762	29378	1
522	1	51	425	18443	1
584	1	51	397	11997	1
1236	1	51	2162	66900	1
535	3	51	609	35000	1

```
In [13]: y_train.head(10)
```

```
Out[13]: 527    9990
129    9500
602    7590
331    8750
323    9100
1358   10900
522   10800
584    9999
1236    8500
535   10500
Name: price, dtype: int64
```



```
In [14]: x_test.head(10)
```

```
Out[14]:
```

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1
1414	1	51	762	39751	1
1089	1	51	882	33160	1
1507	1	51	701	17324	1
970	1	51	701	29000	1
1198	1	51	1155	38000	1

```
In [15]: y_test.head(10)
```

```
Out[15]: 481      7900
76      7900
1502    9400
669     8500
1409    9700
1414    9900
1089    9900
1507    9950
970     10700
1198     8999
Name: price, dtype: int64
```

```
!pip3 install scikit-learn
```

test and train

```
In [16]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [17]: x_test.head(10)
```

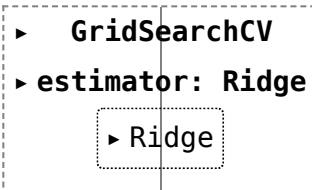
Out[17]:

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1
1414	1	51	762	39751	1
1089	1	51	882	33160	1
1507	1	51	701	17324	1
970	1	51	701	29000	1
1198	1	51	1155	38000	1

Ridge regression

```
In [18]: from sklearn.model_selection import GridSearchCV  
from sklearn.linear_model import Ridge
```

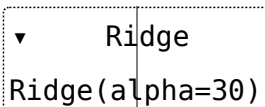
```
In [19]: alpha=[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,1,5,10,20,30]
         ridge=Ridge()
         parameters={'alpha':alpha}
         ridge_regressor=GridSearchCV(ridge,parameters)
         ridge_regressor.fit(x_train,y_train)
```

```
Out[19]: 
```

```
In [20]: ridge_regressor.best_params_
```

```
Out[20]: {'alpha': 30}
```

```
In [21]: ridge=Ridge(alpha=30)
         ridge.fit(x_train,y_train)
```

```
Out[21]: 
```

```
In [22]: y_pred_ridge=ridge.predict(x_test)
         y_pred_ridge
```

```
Out[22]: array([ 5987.68298442,  7272.49041922,  9839.84769665,  9696.77540486,
                10012.04086199,  9628.2868526 ,  9646.94516016, 10090.96059178,
                9877.09434131,  9326.08898223, 10405.82324867,  7711.38762323,
                7676.22360235,  6686.18582795,  9637.23663747, 10345.19967839,
                9401.55433479,  7697.05744783,  4753.93014699, 10426.80946613,
                10340.86050971, 10363.5164179 ,  7639.33873562,  9925.71907421,
                7121.95942411,  9113.15929334,  4919.83330015,  6932.90675142,
                7788.26023667,  9598.07860192,  7313.64648609,  5213.57849239,
                5560.81625893,  5182.12147596,  8921.18592134,  5650.16555828,
                9932.50093896,  8231.7985179 ,  6263.1588039 ,  8570.95039919,
                9747.36874214,  6885.53376523,  8986.75029841, 10319.07112826,
                8669.87879661, 10300.47496653,  9180.81858681,  8843.8523145 ,
                7035.74653423,  9049.29996554,  9389.06587617, 10265.97017653,
                10044.71397804,  6868.062338 ,  9759.8967412 ,  9466.53179916,
                9309.97063726, 10414.69338808,  9768.80893086,  7315.85828673,
                10080.53675451,  6989.45302904,  9827.12942945,  7128.10041765,
                6398.29947219,  9969.70216375,  9754.65880446,  8628.96661395,
                8433.68258496,  6480.90271862,  7746.39132368,  6638.20569737,
                8325.17631535, 10410.63578515,  7336.78464334,  8539.37888515,
                8704.15000000, 10000.00000000,  7350.56101604,  8531.06014240]
```

mean_squared error

```
In [23]: from sklearn.metrics import mean_squared_error#mean_squared error
         Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
         Ridge_Error
```

```
Out[23]: 590569.9121697355
```

finding the efficieny

```
In [24]: from sklearn.metrics import r2_score
         r2_score(y_test,y_pred_ridge)#finding the efficieny
```

```
Out[24]: 0.8391885506165899
```

```
In [25]: a=drop.loc[drop.model==1]
a
```

```
Out[25]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
3	1	51	2739	160000	1	6000
6	1	51	731	11600	1	10750
7	1	51	1521	49076	1	9190
11	1	51	366	17500	1	10990
...
1528	1	51	2861	126000	1	5500
1529	1	51	731	22551	1	9900
1530	1	51	670	29000	1	10800
1534	1	74	3835	112000	1	4600
1536	1	51	2557	80750	1	5990

1094 rows × 6 columns

```
In [26]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
```

```
In [27]: alpha=[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,1,5,10,20,30]
ridge=Ridge()
parameters={'alpha':alpha}
ridge_regressor=GridSearchCV(ridge,parameters)
ridge_regressor.fit(x_train,y_train)
```

```
Out[27]:
```

```

> GridSearchCV
> estimator: Ridge
    > Ridge

```

```
In [28]: ridge_regressor.best_params_
```

```
Out[28]: {'alpha': 30}
```

```
In [29]: ridge=Ridge(alpha=30)
         ridge.fit(x_train,y_train)
         y_pred_ridge=ridge.predict(x_test)
```

```
In [30]: from sklearn.metrics import mean_squared_error#mean_squared error
         Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
         Ridge_Error
```

```
Out[30]: 590569.9121697355
```

```
In [31]: from sklearn.metrics import r2_score
         r2_score(y_test,y_pred_ridge)#finding the efficieny
```

```
Out[31]: 0.8391885506165899
```

```
In [32]: results=pd.DataFrame(columns=['price','predicted'])
results['price']=y_test
results['predicted']=y_pred_ridge
results=results.reset_index()
results['ID']=results.index
results.head(10)
```

Out[32]:

	index	price	predicted	ID
0	481	7900	5987.682984	0
1	76	7900	7272.490419	1
2	1502	9400	9839.847697	2
3	669	8500	9696.775405	3
4	1409	9700	10012.040862	4
5	1414	9900	9628.286853	5
6	1089	9900	9646.945160	6
7	1507	9950	10090.960592	7
8	970	10700	9877.094341	8
9	1198	8999	9326.088982	9

```
In [33]: results['actual price']=results.apply(lambda column:column.price-column.predicted,axis=1)
results['actual price']=results.apply(lambda column:column.price-column.predicted,axis=1)
results
```

Out[33]:


	index	price	predicted	ID	actual price	
	0	481	7900	5987.682984	0	1912.317016
	1	76	7900	7272.490419	1	627.509581
	2	1502	9400	9839.847697	2	-439.847697
	3	669	8500	9696.775405	3	-1196.775405
	4	1409	9700	10012.040862	4	-312.040862

	503	291	10900	10005.311518	503	894.688482
	504	596	5699	6400.852430	504	-701.852430
	505	1489	9500	10096.776914	505	-596.776914
	506	1436	6990	8358.743798	506	-1368.743798
	507	575	10900	10343.148204	507	556.851796

508 rows × 5 columns


```
In [34]: sb.lineplot(x='ID',y='price',data=results.head(50))  
sb.lineplot(x='ID',y='predicted',data=results.head(50))  
mp.plot()
```

Out[34]: []



In []:

In []:

In []: